


Server Internals

Kameswari Chebrolu

Department of CSE, IIT
Bombay



<https://i.pinimg.com/736x/d5/12/48/d5124806ed216c46103e65c475d3bb81.jpg>

Background: Overall Outline

- ~~What constitutes a webpage?~~
- ~~What goes on inside a Browser?~~
 - ~~What standard security mechanisms implemented?~~
- ~~How do client and server communicate?~~
 - ~~HTTP/HTTPS protocol~~
 - ~~Session Management via cookies and tokens~~
- **How does a web server process requests and generate responses?**
 - **Static vs Dynamic content**

Server Content

- Servers serve two types of content: Static and Dynamic
- Focus: Static (for now)
 - HTML files, image files, javascript files
 - Locate the requested file on file system and return unaltered in HTTP response

Server-side

Client-side



Files

Pre-created:
HTML
CSS
JS and other files



Web Server

HTTP Request

HTTP Response



Browser

- Modern web servers may do some things extra!
 - Dynamically compress large resource files using gzip
 - Reduces bandwidth used in the response
 - Add caching headers in HTTP responses to instruct browser to cache
 - Decreases page load times in future if same file needed!

Web Server Software

- Plays a key role in managing connections and serving static (as well as dynamic content)
- Apache HTTP Server (Apache): Oldest and most widely used open-source web server
 - Process driven architecture (not very memory efficient)
 - Configurability via files
 - main configuration file (apache2.conf)
 - .htaccess files in specific directories (override global settings)



- Rich ecosystem of modules that extend functionality
 - Modules handle various tasks such as authentication, URL rewriting, caching, load balancing, logging etc
 - mod_ssl: Adds support for SSL/TLS encryption.
 - mod_rewrite: Enables powerful URL rewriting and redirection
 - <http://example.com/product/123> rewritten to <http://example.com/product.php?id=123>
 - mod_proxy: Provides support for proxying requests to other servers.
 - mod_auth: Handles authentication and access control

- Nginx: (pronounced "engine-x")
 - Configuration typically done via `nginx.conf`
 - Lightweight, high-performance web server
 - Event-driven architecture → handles concurrency very well
 - Also supports reverse proxying
 - Nginx often preferred for static content and as reverse proxy
 - Focuses on a small core feature set
 - Number of modules not as extensive as Apache's



- Microsoft Internet Information Services (IIS):
Supports various web technologies and integrates well with other Microsoft products



Directory Structure (Linux)

- Specifics vary based on distribution and server software
- Web Server Root: Base directory from which web server serves files (user requested)
 - HTML, CSS , Javascript files are placed here
 - For Apache, the default root directory is often `/var/www/html`
 - For Nginx, the default root directory is `/usr/share/nginx/html`
- Server Configuration Files: Configuration files for a web server are usually found in `/etc`
 - Apache configuration files might be in `/etc/apache2`
 - Nginx configuration files might be in `/etc/nginx`

- Logs: access logs and error logs, are typically found in /var/log
 - For Apache, logs might be in /var/log/apache2
 - For Nginx, logs might be in /var/log/nginx
- Virtual Hosts (if applicable): One web server can host multiple websites
 - Configurations for these virtual hosts may be stored in a subdirectory of the server's configuration directory
 - E.g., /etc/apache2/sites-available/ for Apache

Sample Apache Configuration

Global configuration

ServerRoot "/etc/apache2"

Listen on port 80

Listen 80

Server-wide defaults

<Directory />

Options FollowSymLinks

AllowOverride None

Require all denied

</Directory>

<Directory /var/www/>

Options Indexes FollowSymLinks

AllowOverride None

Require all granted

</Directory>

Logging

ErrorLog \${APACHE_LOG_DIR}/error.log

LogLevel warn

CustomLog \${APACHE_LOG_DIR}/access.log combined

Include module configurations

IncludeOptional mods-enabled/*.load

IncludeOptional mods-enabled/*.conf

Include additional directory configurations

IncludeOptional conf-enabled/*.conf

Virtual hosts

IncludeOptional sites-enabled/*.conf

Explanation

- **ServerRoot Directive:** top-level directory where server's configuration files and assets are located
 - Set to `"/etc/apache2"`
- **Listen Directive:** port on which Apache listens for incoming connections
 - Set to default port 80

- **<Directory> Directives**: Define settings for specific directories
 - **First <Directory> block sets options for the root directory ("/")**
 - Second block for the default document root ("/var/www/")
 - Options: Specifies the default behavior for the specified directory
 - "FollowSymLinks" option allows symbolic links (symlinks) to be followed
 - **Indexes" option allows server to generate and display directory listings for directories that don't have an index file (e.g. index.html)**

- AllowOverride: Which directives can be overridden using .htaccess files
 - All: Allows all directives to be overridden by .htaccess files
 - None: Don't allow directives to be overridden by .htaccess files
- Require: Access control rules
 - all: It refers to all users or clients, effectively meaning "everyone."
 - Require ip 192.168.1.0/24: Restricts access to only clients within the specified IP range
 - access is denied by default for / and explicitly allowed for document root

- Logging Section: Configures error and access logs
 - `${APACHE_LOG_DIR}` is a variable often set to `var/log/apache2`
 - Defined in `/etc/apache2/envvars`
 - `ErrorLog`: Specifies the path to the error log file
 - `LogLevel`: Sets the verbosity level
 - `CustomLog`: Defines the format and location of the access log
 - “Combined” is one of a predefined log format

- Include Directives: Includes additional files
 - The .load files typically contain LoadModule directives, which load Apache modules.
 - The .conf files typically contain configuration settings
 - Related to the loaded modules (under mods-enabled)
 - Or other configurations (conf-enabled or sites-enabled for virtual hosting etc)

Motivation: Dynamic Content

- Add new item to inventory → painful to manually create a new product page
 - Repeat lot of code across each page
 - Any change to page structure → changes to many pages!
- Good to have an automated way to read from database and create product page!
- Other examples:
 - Display score/marks specific to logged in student
 - Process form submissions at server and provide feedback to user

Dynamic Content

- Content dynamically generated on the fly in response to user request
- Example:
 - Receive a HTTP GET Request for a product
 - Server determines the product ID
 - Fetches data from database
 - Constructs the HTML page for the response by inserting the data into a HTML template
 - Any changes need to be done in one place, in a single template, and not across many static pages!

Technology

- Server-Side Scripting Languages: Executed on server to create dynamic content
 - **PHP**, Python (Django), Ruby (Ruby on Rails), Node.js etc
- **Templates**: Help separate logic from presentation
 - Servers embed dynamic content within HTML templates
- Databases: Dynamic content generated via data stored and retrieved from databases
 - MySQL, PostgreSQL, MongoDB etc
 - Server-side scripts use queries (e.g. SQL) to interact with databases

- Web Frameworks: Frameworks like **Flask**, Django, Ruby on Rails, and Express.js provide structured web application development
 - Often based on Model-View-Controller (MVC)
 - Separate application into
 - models (data and business logic)
 - views (presentation)
 - controllers (handling user input and interaction)

- **Web APIs** (Application Programming Interfaces):
Dynamic content can be sourced from external APIs
 - Websites can pull data from third-party services or other platforms
 - RESTful APIs are commonly used
- Real-Time Technologies:
 - **WebSockets**: provide a full-duplex communication channel between server and client
 - Server-Sent Events (SSE): a mechanism for sending updates from the server to client

Server-side Scripting: PHP

- PHP stands for Hypertext Preprocessor
- Embedded within HTML code itself
 - Allows developers to seamlessly mix server-side logic with HTML markup
 - Executed on the server, and the resulting HTML is sent to client



Example

<http://www.example.com/welcome.php?name=Chotu>

(question mark (?)) indicates the beginning of a query string. Query string is a set of key-value pairs separated by ampersands (&))

- `isset($_GET['name'])`: Checks if 'name' parameter is present in URL
- `$_GET['name']`: retrieve user's name
- `htmlspecialchars()` sanitizes the input
- Displays a personalized greeting using the user's name
- If 'name' parameter not provided, another message is shown

```
<body>
  <h1>PHP Example</h1>

  <?php
    // Check if the 'name' parameter is present in the URL
    if (isset($_GET['name'])) {
      // Get the user's name from the 'name' parameter
      $userName = htmlspecialchars($_GET['name']);

      // Display a personalized greeting
      echo "<p>Welcome, {$userName}!</p>";
    } else {
      // Display a message if 'name' parameter is not provided
      echo "<p>Please provide your name in the URL.</p>";
    }
  ?>
</body>
```


Workflow with PHP

1. **Client Request:** A user enters URL or clicks on a link
2. **Web Server Handling:** Web server (e.g., Apache or Nginx) receives the request
3. **PHP Processing:**
 - Requested resource is a PHP file → web server hands over request to PHP interpreter
 - PHP scripts often interact with databases (e.g., MySQL, PostgreSQL) to retrieve or update data.
 - PHP generates dynamic content, typically HTML, based on the logic in PHP script and any data retrieved from database
4. **Response to Client:**
 - Generated content is sent back as an HTTP response to user's web browser

In file `/etc/apache2/apache2.conf` include php related configuration

```
LoadModule php_module modules/libphp.so
```

```
AddHandler application/x-httpd-php .php
```

(LoadModule loads php module from given path and names it `php_module`)

AddHandler associates given MIME type to files ending in `.php`

`php_module` communicates with the PHP interpreter to execute the PHP script

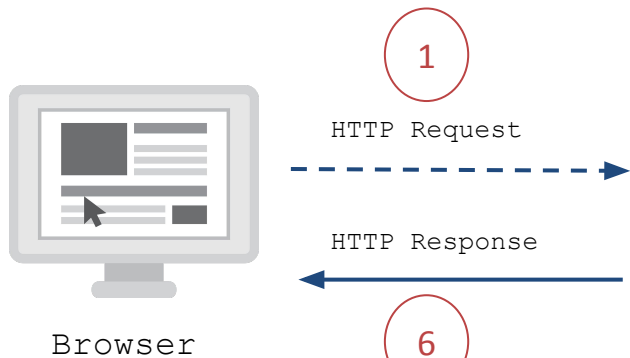
PHP interpreter processes the script, executes the PHP code and generates dynamic content (HTML) and sends it back to `php_module`

module sends the HTTP response to client

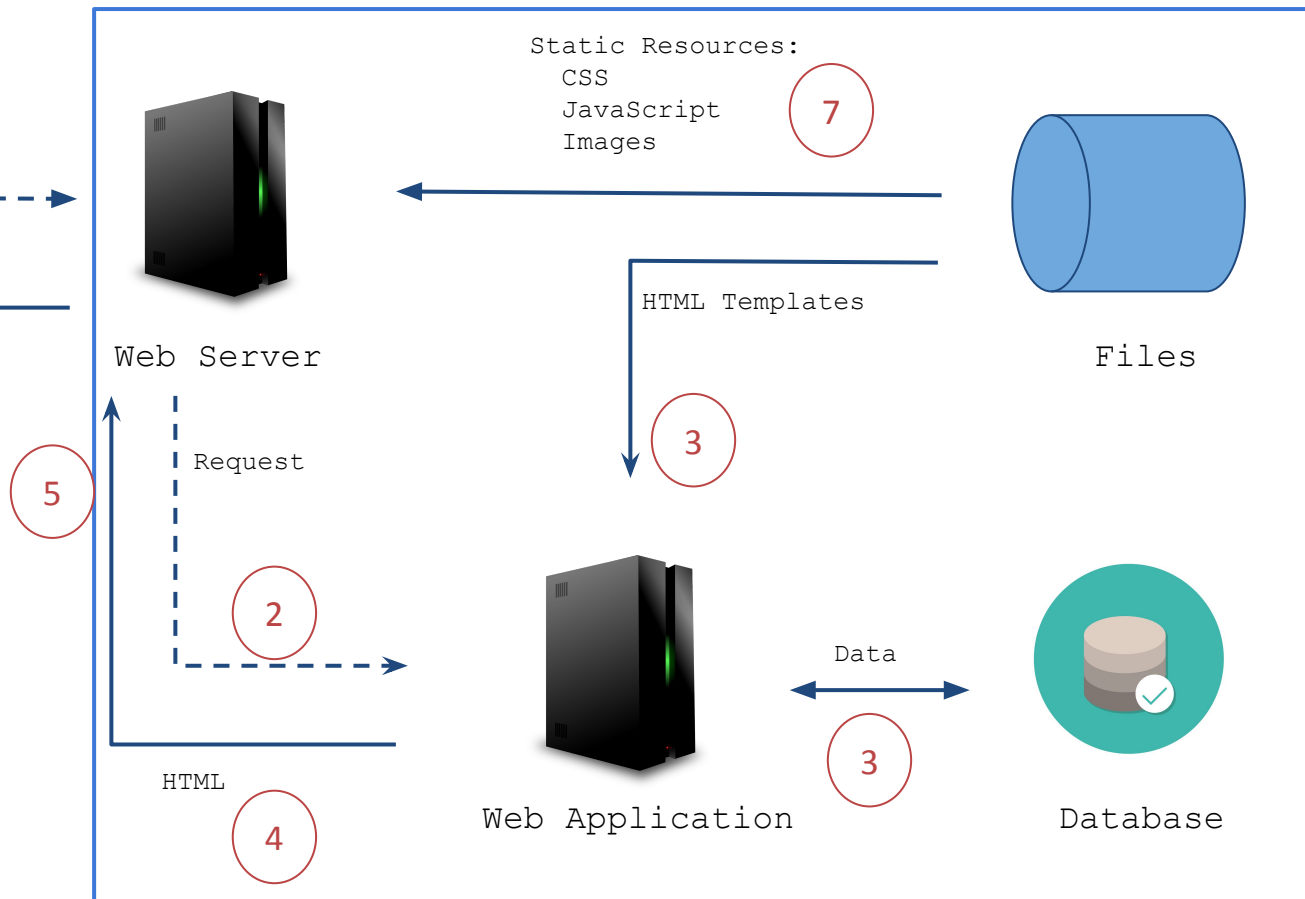
Web Application Frameworks

- Software frameworks make it easier to write, maintain and scale web applications
 - Provide tons of tools and libraries that simplify common tasks
 - Routing URLs to appropriate handlers, template engines, interacting with databases, supporting sessions and user authorization, formatting output (e.g. HTML, JSON, XML), and improving security etc
- Examples of popular web frameworks
 - Django (Python), Ruby on Rails (Ruby), Flask (Python), Spring Boot (Java)

Client-Side



Server-Side



Workflow

- Web browser sends a HTTP GET request
 - E.g. <https://www.example.com/product/123> or <https://www.example.com/product?id=123>
 - Note: GET request is used because request is only fetching data
- Web Server detects request is "dynamic" and forwards it to Web Application for processing
 - How does it know where to forward?
 - Based on pattern matching rules defined in its configuration

- Web Application identifies intention of request is to get product information of product with id=123
 - Gets required information from the database
 - Dynamically creates a HTML page by putting retrieved data into placeholders inside a **HTML template**
- Returns generated HTML to web browser (via the Web Server), along with HTTP status code (200 for success)
 - Web Application may return another code is something fails
 - E.g. "404" to indicate product does not exist

- Web Browser processes returned HTML
 - Can send separate requests for other CSS or JavaScript files
 - Web Server will return such static files directly
 - Correct file handling based on configuration rules and URL pattern matching!
- Note: Server-side code does not have to return HTML files only, it can dynamically create and return other types of files
 - E.g. text, PDF, CSV, JSON, XML, etc

Example: Flask

```
from flask import Flask, render_template, request
```

```
app = Flask(__name__)
```

```
# Route to the home page
```

```
@app.route('/')
```

```
def home():
```

```
# Get the 'username' query parameter from the URL
```

```
username = request.args.get('username', 'Guest')
```

```
# Data to be passed to the template
```

```
user = {'username': username}
```

```
# Render the template 'index.html' with the provided data
```

```
return render_template('index.html', user=user)
```

```
if __name__ == '__main__':
```

```
    app.run(debug=True)
```

- Flask application defines a route for the home page ('/')
 - When a user accesses this page, the `home` function is called
- request object from Flask is used to access the query parameters in the URL
 - retrieves 'username' parameter, and if it's not present, it defaults to 'Guest'
- A dictionary (`user`) is created, and the `render_template` function is called to render 'index.html'
 - `user` data is passed to it
- In the template (`index.html`), Jinja2 syntax (`{{ user.username }}`) is used to insert dynamic data into the HTML content

(Jinja2 is a template engine for Python programming language)

Template (index.html)

```
<!DOCTYPE html>
<html lang="en">
<head>
  <title>Flask Template Example</title>
</head>
<body>
  <h1>Hello, {{ user.username }}!</h1>
  <p>This is a simple example of using templates with Flask.</p>
</body>
</html>
```


Web APIs

- API: Application Programming Interface
 - Set of rules and protocols that allows one application to interact with another
- Web API?
 - Enables communication and data exchange between web servers and also client-server
 - Typically use HTTP (Hypertext Transfer Protocol) as the communication protocol

- Many web APIs follow REST principles
 - RESTful (Representational State Transfer) APIs
 - **Stateless**, scalable, often idempotent and adhere to standard HTTP methods (GET, POST, PUT, DELETE) for performing operations on resources
 - Stateless: each request from a client contains all info needed for server to respond
 - No client state stored between requests.

- Web APIs expose specific URLs (endpoints) to which clients send HTTP requests
 - Each endpoint represents a specific resource or functionality provided by the API
- Web APIs exchange data in standardized formats, such as JSON (JavaScript Object Notation) or XML (eXtensible Markup Language)

Example

Endpoint: <http://api.example.com/users> (This URL serves as the base for all user-related operations in the API)

GET /users HTTP/1.1

Host: api.example.com

Response:

HTTP/1.1 200 OK

Content-Type: application/json

```
[
  {
    "id": 1,
    "username": "john_doe",
    "email": "john.doe@example.com"
  },
  {
    "id": 2,
    "username": "jane_smith",
    "email": "jane.smith@example.com"
  }
]
```

GET /users/1 HTTP/1.1
Host: api.example.com

Response:

HTTP/1.1 200 OK
Content-Type: application/json

```
{  
  "id": 1,  
  "username": "john_doe",  
  "email": "john.doe@example.com"  
}
```

POST /users HTTP/1.1
Host: api.example.com
Content-Type: application/json

```
{  
  "username": "new_user",  
  "email": "new.user@example.com"  
}
```

Response
HTTP/1.1 201 Created
Content-Type: application/json

```
{  
  "id": 3,  
  "username": "new_user",  
  "email": "new.user@example.com"  
}
```

- For security,
 - Web APIs often require authentication
 - Clients must provide valid credentials (e.g., API keys, tokens) in the requests
 - Authorization also important
 - What actions or resources a particular user or application can access?
 - OAuth2.0 tokens and JWTs contain relevant info (user roles/permissions)

- Cross-Origin Resource Sharing (CORS) help web pages hosted on different domains to access web APIs
 - Relevant CORS headers specify which domains are allowed to make requests!
 - More in depth coverage later under attacks!
- Popular third-party web APIs: Twitter API, Google Maps API, GitHub API
 - Developers use these APIs to integrate third-party services into their applications

- As an aside: there are Browser APIs also
 - Extend functionality of the browser
 - E.g. Geolocation API returns coordinates of where browser is located
 - Web form API for input data validation
 - Web storage api for storing and retrieving data within browser
 - Web fetch api (saw as part of AJAX)s

Data

- Exchange of data between a client and a server can occur in various formats
 - Already saw: html, css, js files; images in form of png, jpg; pdfs, docs etc
- APIs often involve transfer of structured data
 - JSON and XML are commonly used
 - Earlier in web's history, XML was in vogue, now JSON is more used!

JavaScript Object Notation (JSON)

- A lightweight data interchange format
 - Easy for humans to read/write
 - Also for machines to parse and generate
- Data represented as key-value pairs
 - Supports data types such as strings, numbers, booleans, arrays, and objects
 - Data is enclosed in curly braces {}
 - Key-value pairs are separated by colons
 - Arrays are represented using square brackets []
 - Case-sensitive and whitespace-insensitive

```
{  
  "name": "Ravi",  
  "age": 22,  
  "isStudent": true,  
  "grades": [90, 85, 92],  
  "address": {  
    "city": "Mumbai",  
    "zipcode": "400076"  
  }  
}
```

- MIME type: application/json
- Limitations:
 - no support for comments
 - no native support for binary data
 - Inability to represent circular references in objects
- More advanced features needed: use XML, Protocol Buffers, or GraphQL

Web Sockets

- WebSockets provide a full-duplex communication channel over a single, long-lived connection
 - Full-Duplex Communication: data can be sent in both directions simultaneously
 - Traditional HTTP follows a half-duplex model; client sends request, server responds!
 - Persistent Connection: connection remains open for as long as both client and server want

- WebSocket Protocol operates over a single, dedicated TCP connection.
 - Need to use `ws://` and `wss://`
- Suitable for real-Time Applications
 - E.g. chat, online gaming, financial trading, live updates etc
- More details later under attacks!

Summary as History

- Static Web Pages (1990s): No interactivity and dynamic content
- JavaScript (1995): interactivity within browser
- AJAX (early 2000s): use JavaScript to make asynchronous requests
 - For more responsive and interactive web applications
- Server-Side Technologies (PHP, ASP, JSP, early 2000s): server-side scripting languages
 - Embed server-side logic into web pages to generate dynamic content

- Web Frameworks and MVC Architecture (mid-2000s): Model-View-Controller (MVC) style architectural patterns
 - Organize code, separating concerns and promoting development of dynamic and maintainable web applications
- Web 2.0 and Rich Internet Applications (RIAs, mid-2000s): shift towards more user-centric and interactive web experiences
 - Based on Adobe Flash, advanced JavaScript libraries (e.g., jQuery)

- Single Page Applications (SPAs, 2010s): load a single HTML page and dynamically update content
 - Based on Angular, React, and Vue.js etc
- WebAssembly (Wasm, 2015s): a binary instruction format that enables high-performance execution of code (C, C++, Rust) on web browsers
- Progressive Web Apps (PWAs, current): Use modern web capabilities to deliver an app-like experience across different devices
 - Offline support, push notifications, and improved interactivity and performance

References

- https://developer.mozilla.org/en-US/docs/Learn/Server-side/First_steps/Introduction
- <https://www.w3schools.com/php/>
- <https://www.geeksforgeeks.org/flask-tutorial/> and <https://www.geeksforgeeks.org/getting-started-with-jinja-template/>
- <https://www.tutorialspoint.com/restful/index.htm> (REST api)
- <https://www.tutorialspoint.com/websockets/index.htm>
- https://www.w3schools.com/js/js_json_intro.asp (json)